

CSCI 4061: Processes and Environment

Chris Kauffman

*Last Updated:
Wed Jan 30 19:02:22 CST 2019*

Logistics

Reading

Stevens and Rago, Ch 7-8

Goals Today

- ▶ Finish Basics Overview
- ▶ Process Lifecycle
- ▶ Killing programs
- ▶ Process memory layout

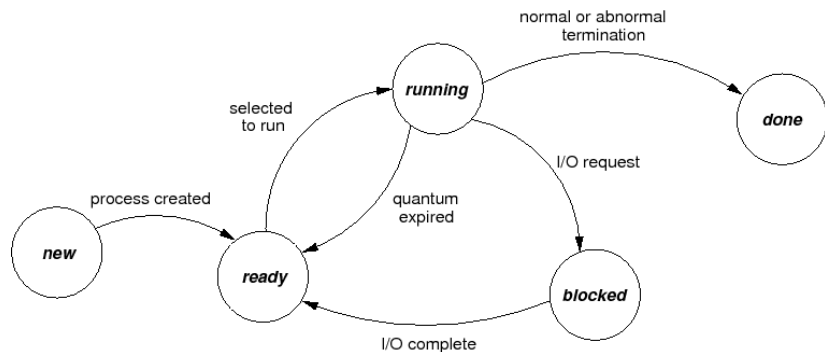
Lab01: Makefiles / Compilation

- ▶ Due Monday night
- ▶ Feedback?

Reminder: AGREEMENT

Download `AGREEMENT.txt` from the schedule page, sign and submit to Canvas

Process Life Cycle



Source: Saverio Perugini, lecture notes

ps and top show running process status

These shell commands show a STAT or S columns corresponding loosely to process states.

STAT	Meaning
<i>Common</i>	
R	running or runnable (on run queue)
S	interruptible sleep (waiting for an event to complete)
T	stopped, either by a job control signal or being traced.
Z	defunct ("zombie") process, terminated but not reaped by parent.
<i>Less Common</i>	
D	uninterruptible sleep (usually IO)
W	paging (not valid since the 2.6.xx kernel)
X	dead (should never be seen)


Source: man page for ps

Will talk more about zombies and orphans later



States of a Living Process

- ▶ Note inclusion of Kernel/OS here
- ▶ **Interrupt and Sys Calls** start running code in the operating system
- ▶ Interrupt/Signal can come from software or hardware
- ▶ **Context switch** starts running another process, only happens when one process is safely tucked in and put to **sleep**



Source: *Design of the Unix Operating System* by Maurice Bach

Terminal: Foreground/Background Processes

- ▶ Type a program into the terminal, press enter
- ▶ Starts a process in the **foreground** of the terminal
- ▶ Input from user typing, output to terminal screen
- ▶ **Suspend** foreground process with `Ctrl-z`, gets prompt back
- ▶ **Terminate** foreground process (usually) with `Ctrl-c`
- ▶ Run a process in the background
 - ▶ Start with `&` at end: `ls &`
 - ▶ Move job 2 to bg: `bg %2`
- ▶ Move job 2 to foreground: `fg %4`
- ▶ Terminals running programs with jobs

Murdering Processes

Question: What command ends processes that are misbehaving?

Keystrokes to Remember

Ctrl-c Send the interrupt signal, kills most processes

Ctrl-z Send the stop signal, puts process to sleep

Easy to Kill

- ▶ yes spits output to the screen continuously
- ▶ End it from the terminal its started in
- ▶ Suspend it then, end it
- ▶ Kill it from a different terminal

Harder to Kill

- ▶ Consider the program `no_interrupts.c`
- ▶ Ignores some common signals
- ▶ Need to use the big stick for this one:
`kill -9 1234` OR
`pkill -9 a.out`

Exercise: Basic Job Control

Give a sequence of commands / keystrokes to...

Misbehaving

- ▶ Compile `no_interrupts.c` to a program named `invincible`
- ▶ Run `invincible`
- ▶ Try to end the process by sending it the interrupt signal
- ▶ In a separate terminal, end the `invincible` program

Edit / Build Seq

- ▶ Edit `Makefile` with `vi`
- ▶ Suspend `vi` (don't quit it)
- ▶ Run the `Makefile`
- ▶ Terminate before completing build
- ▶ Bring back `vi` to continue editing

Recall: Program Memory

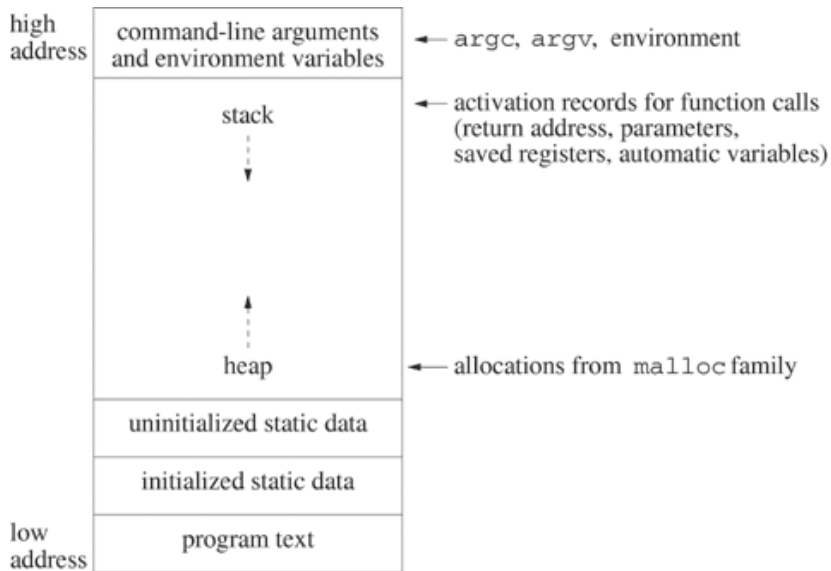
- ▶ What are the 4 memory areas to a C program we've discussed OR that you know from previous courses?
- ▶ Give an example of how one creates variables/values in each area of memory

Answers: Program Memory

- ▶ What are the 4 memory areas to a C program we've discussed OR that you know from previous courses?
 1. Stack: automatic, push/pop with function calls
 2. Heap: malloc() and free()
 3. Global: variables outside functions, static vars
 4. Text: Assembly instructions
- ▶ Give an example of how one creates variables/values in each area of memory

```
1 #include <stdlib.h>
2 int glob1 = 2;           // global var
3 int func(int *a){        // param stack var
4     int b = 2 * (*a);    // local stack var
5     return b;            // de-allocate locals in func()
6 }
7 int main(){              // main entry point
8     int x = 5;           // local stack var
9     int c = func(&x);     // local stack var
10    int *p = malloc(sizeof(int)); // local stack var that points into heap
11    *p = 10;              // modify heap memory
12    glob1 = func(p);      // allocate func() locals and run code
13    free(p);              // deallocate heap mem pointed to p
14    return 0;             // deallocate locals in main()
15 }
16 // all executable code is in the .text memory area as assembly instructions
```

More Detailed Process Memory

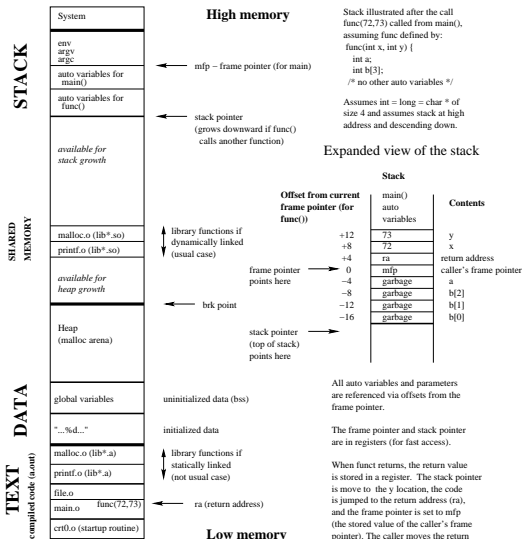


Source: *Unix Systems Programming, Robbins & Robbins*

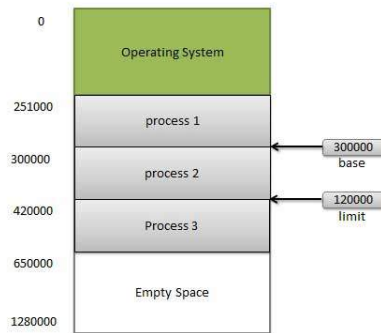
Yet *more* detailed view (Link)

A detailed picture of the virtual memory image, by [Wolf Holzman](#)

Memory Layout (Virtual address space of a C process)



Unix Processes In Memory



Source: Tutorials Point

- ▶ Separate Memory Image for Each Process
- ▶ OS + Hardware keeps processes inside their own address space
- ▶ This is a gross simplification but will suffice until later when we discuss **virtual memory** vs physical memory

This picture should bother you

- ▶ Consequence for program dynamic memory allocation?
- ▶ Problems with running system calls?

Exercise: Memory Problems in C Programs

What you're up against

- ▶ Stack problems: References to stack variables that go away
- ▶ Segmentation Faults: Access memory out of bounds for whole program, via heap or via stack
- ▶ Null pointers dereference: Often results in a segfault as NULL translates to 0x0000 which is off limits
- ▶ Use of uninitialized: variables don't have values by default, assign or get something random
- ▶ Memory Leaks: `malloc()` memory that is not used but never `free()`'d, program gobbles more and more memory
- ▶ Examine results of running `overflow.c`, **EXPLAIN OUTPUT**

Solutions

- ▶ Don't program in C
- ▶ Use a tool to help identify and fix problems
- ▶ Memory Tools on Windows: [Discussion here](#). Synopsis → \$\$\$
- ▶ Memory Tools on Linux/Mac: **Valgrind** → FREE

Valgrind: Memory Tool on Linux/Mac



Valgrind¹ has Memcheck

- ▶ Catches most memory errors²
 - ▶ Use of uninitialized memory
 - ▶ Reading/writing memory after it has been free'd
 - ▶ Reading/writing off the end of malloc'd blocks
 - ▶ Memory leaks
- ▶ Source line of problem happened (but not cause)
- ▶ Super easy to use, installed on lab machines
- ▶ Slows execution of program way down

¹<http://valgrind.org/>

²<http://en.wikipedia.org/wiki/Valgrind>

Valgrind in Action

See some common problems in `badmemory.c`

Debuggers


- ▶ There comes a day when `printf` just isn't enough
- ▶ On that day you will start compiling with `-g` to turn on the debugger
- ▶ Then you will run `gdb myprog`, set some breakpoints, and get to the root of the problem
- ▶ More on debugger later in a lecture/lab

Command Line Arguments

```
int main(int argc, char *argv[])
```

2-arg version of `main()` will be set up to have number of arguments and array of strings in it by whatever started it

```
> cat print13.c
#include <stdio.h>
int main(int argc, char *argv[]){
    printf("%s\n",argv[1]);
    printf("%s\n",argv[3]);
}
> gcc print13.c
> ./a.out -c 10 2.0
-c
2.0
```



`argc` is 4 in this case

Exercise: Print Args

Write a quick C program which prints all of its `argv` elements as strings. There will be `argc` elements in this array.

Environment Variables

All programs can access **environment** variables, name/value pairs used to communicate and alter behavior.

Shell show/set variables

Done with `echo $VARIABLE`

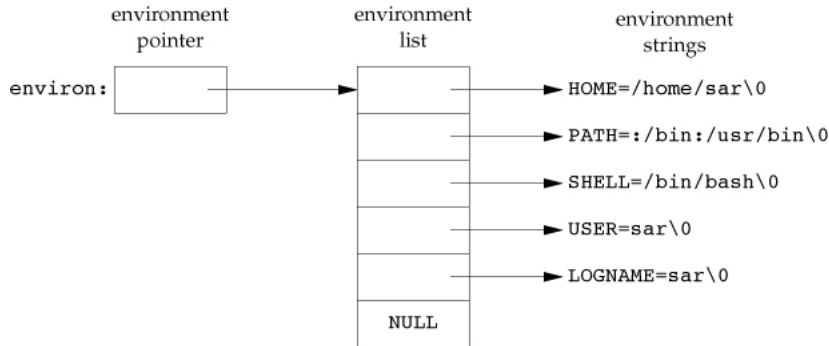
```
> echo $PAGER
less
> PAGER=cat
> echo $PAGER
cat
> echo $PS1
>
> PS1='wicked$ '
wicked$
```

Shell env

Show *all* environment

```
> env
JAVA8_HOME=/usr/lib/jvm/java-8-openjdk
PAGER=less
PWD=/home/kauffman/4061-F2017/lectures/03
HOME=/home/kauffman
BROWSER=chromium
COLUMNS=79
MAIL=/var/spool/mail/kauffman
MANPATH=:/home/kauffman/local/man:/home/k
PATH=/usr/local/sbin:/usr/local/bin:/usr/
PS1=>
...
```

C Programs and Environment Vars



- ▶ Global variable `char **environ` provides array of environment variables in form `VARNAME=VALUE`, null terminated
- ▶ Easier to use the library functions to check/change environment

C Library for Environment Vars

The C Library Provides standard library functions for manipulating environment variables.

```
#include <stdlib.h>

char *getenv(const char *name);
// returns pointer to value associated with name, NULL if not found

int setenv(const char *name, const char *value, int rewrite);
// sets name to value. If name already exists in the environment, then
// (a) if rewrite is nonzero, the existing definition for name is
// first removed; or (b) if rewrite is 0, an existing definition for
// name is not removed, name is not set to the new value, and no error
// occurs. return: 0 if OK, -1 on error

int unsetenv(const char *name);
// removes any definition of name. It is not an error if such a
// definition does not exist. return: 0 if OK, -1 on error

int putenv(char *str);
// str is of form NAME=VALUE, alters environment accordingly. If name
// already exists, its old definition is first removed. Don't use with
// stack strings. Returns: 0 if OK, nonzero on error.
```

Exercise: Manipulate Environment Vars

Write a short C program which behaves as indicated in the demo

- ▶ Prints ROCK and VOLUME environment variables
- ▶ If ROCK is set to anything, change VOLUME to "11"

Use these functions

```
char *getenv(const char *name);  
// NULL if name not set  
// otherwise pointer to value  
  
int setenv(const char *name,  
           const char *value,  
           int rewrite);  
// Change name value pair,  
// if rewrite is 1,  
// overwrite previous definitions
```

Note the use of export to ensure child processes see the environment variables

```
> unset ROCK  
> unset VOLUME  
> gcc environment_vars.c  
> a.out  
ROCK not set  
VOLUME is not set  
> export VOLUME=7  
> a.out  
ROCK not set  
VOLUME is 7  
> export ROCK=yes  
> a.out  
ROCK is yes  
Turning VOLUME to 11  
VOLUME is 11  
> echo $VOLUME  
7
```

Note also that the program does not change the shell's values for ROCK: no child can change a parent's values (or mind)

Answers: Manipulate Environment Vars

See 03-process-basics-code/environment_vars.c

```
1 // environment_vars.c: solution to in-class exercise showing how to
2 // check and set environment variables via the standard getenv() and
3 // setenv() functions.
4 #include <stdlib.h>
5 #include <stdio.h>
6
7 int main(int argc, char *argv[]){
8
9     char *verbosity = getenv("ROCK");
10    if(verbosity == NULL){
11        printf("ROCK not set\n");
12    }
13    else{
14        printf("ROCK is %s\n",verbosity);
15        printf("Turning VOLUME to 11\n");
16        int fail = setenv("VOLUME","11",1);
17        if(fail){
18            printf("Couldn't change VOLUME\n");
19        }
20    }
21    char *volume = getenv("VOLUME");
22    if(volume == NULL){
23        volume = "not set";
24    }
25    printf("VOLUME is %s\n",volume);
26    return 0;
27 }
```