

CSCI 4061: Making Processes

Chris Kauffman

*Last Updated:
Wed Jan 30 18:37:33 CST 2019*

Logistics

Reading

- ▶ Robbins and Robbins, Ch 3
- ▶ OR Stevens and Rago, Ch 8

Goals

- ▶ Project 1
- ▶ Environment Variables
- ▶ Creating Child Processes
- ▶ Waiting for them
- ▶ Running other programs

Lab02: `fork()`, `wait()`, `exec()`

- ▶ All things you'll need in first project
- ▶ Feedback on content
- ▶ Feedback on **grading policy**

Project 1

- ▶ Spec will go up later today
- ▶ Due in about 2.5 weeks
- ▶ Groups of 1 or 2

Overview of Process Creation/Coordination

`getpid() / getppid()`

- ▶ Get process ID of the currently running process
- ▶ Get parent process ID

`fork()`

- ▶ Create a child process
- ▶ Identical to parent EXCEPT for return value of `fork()` call
- ▶ Determines child/parent

`wait() / waitpid()`

- ▶ Wait for any child to finish (`wait`)
- ▶ Wait for a specific child to finish (`waitpid`)
- ▶ Get return status of child

`exec() family`

- ▶ Replace currently running process with a different image
- ▶ Process becomes something else losing previous code
- ▶ Focus on `execvp()`

Overview of Process Creation/Coordination

getpid() / getppid()

```
pid_t my_pid = getpid();
printf("I'm proces %d\n",my_pid);
pid_t par_pid = getppid();
printf("My parent is %d\n",par_pid);
```

fork()

```
pid_t child_pid = fork();
if(child_pid == 0){
    printf("Child!\n");
}
else{
    printf("Parent!\n");
}
```

wait() / waitpid()

```
int status;
waitpid(child_pid, &status, 0);
printf("Child %d done, status %d\n",
       child_pid, status);
```

exec() family

```
char *new_argv[] = {"ls","-l",NULL};
char *command = "ls";
printf("Goodbye old code, hello LS!\n");
execvp(command, new_argv);
```

Exercise: Standard Use: Get Child to Do Something

Child Labor

- ▶ Examine the file `child_labor.c` and discuss
- ▶ Makes use of `getpid()`, `getppid()`, `fork()`, `execvp()`

Child Waiting

- ▶ `child_labor.c` has concurrency issues: parent/child output mixed
- ▶ **Modify** with a call to `wait()` to ensure parent output comes AFTER child output

Answers: Standard Use: Get Child to Do Something

```
1 // child_labor.c: demonstrate the basics of fork/exec to launch a
2 // child process to do "labor"; e.g. run a another program via exec.
3 #include <stdio.h>
4 #include <stdlib.h>
5 #include <sys/wait.h>
6 #include <unistd.h>
7
8 int main(int argc, char* argv){
9
10     char *child_argv[] = {"ls", "-l", "-ah", NULL};           // argument array to child, must end with NULL
11     char *child_cmd = "ls";                                   // actual command to run, must be on path
12
13     printf("I'm %d, and I really don't feel like '%s'ing\n",
14           getpid(), child_cmd);                                // use of getpid() to get current PID
15     printf("I have a solution\n");
16
17     pid_t child_pid = fork();                                  // clone a child
18
19     if(child_pid == 0){                                        // child will have a 0 here
20         printf(" I'm %d My pa '%d' wants me to '%s'. This sucks.\n",
21               getpid(), getppid(), child_cmd);                // use of getpid() and getppid()
22
23         execvp(child_cmd, child_argv);                         // replace running image with child_cmd
24
25         printf(" I don't feel like myself anymore...\n");      // unreachable statement
26     }
27     else{                                                      // parent will see nonzero in child_pid
28         printf("Great, junior %d is taking care of that\n",
29               child_pid);
30     }
31     return 0;
32 }
```

Answers: Standard Use: Get Child to Do Something

```
1 // child_wait.c: fork/exec plus parent waits for child to
2 // complete printing after each time.
3
4 #include <stdio.h>
5 #include <stdlib.h>
6 #include <sys/wait.h>
7 #include <unistd.h>
8
9 int main(int argc, char* argv){
10
11     char *child_argv[] = {"ls", "-l", "-ah", NULL};           // argument array to child, must end with NULL
12     char *child_cmd = "ls";                                   // actual command to run, must be on path
13
14     // char *child_argv[] = {"/complain", NULL};              // alternative commands
15     // char *child_cmd = "complain";
16
17     printf("I'm %d, and I really don't feel like '%s'ing\n",
18           getpid(), child_cmd);
19     printf("I have a solution\n");
20
21     pid_t child_pid = fork();
22
23     if(child_pid == 0){
24         printf(" I'm %d My pa '%d' wants me to '%s'. This sucks.\n",
25               getpid(), getppid(), child_cmd);
26         execvp(child_cmd, child_argv);
27         printf(" I don't feel like myself anymore...\n"); // unreachable
28     }
29     else{
30         int status;
31         wait(&status);                                     // wait for child to finish, collect status
32         printf("Great, junior %d is done with that '%s'ing\n",
33               child_pid, child_cmd);
34     }
35     return 0;
36 }
```

Exercise: Child Exit Status

- ▶ A successful call to `wait()` sets a status variable giving info about child

```
int status;  
wait(&status);
```

- ▶ Several macros are used to parse out this variable

```
// determine if child actually exited  
// other things like signals can cause  
// wait to return  
if(WIFEXITED(status)){  
  
    // get the return value of program  
    int retval = WEXITSTATUS(status);  
}
```

- ▶ **Modify** `child_labor.c` so that parent checks child exit status
- ▶ Convention: 0 normal, nonzero error, print something if non-zero

```
# program that returns non-zero  
> gcc -o complain complain.c  
  
# EDIT FILE TO HAVE CHILD RUN 'complain'  
> gcc child_labor_wait_returnval.c  
> ./a.out  
I'm 2239, and I really don't feel  
like 'complain'ing  
I have a solution  
    I'm 2240 My pa '2239' wants me to 'complain'.  
    This sucks.  
COMPLAIN: God this sucks. On a scale of 0 to 10  
    I hate pa ...  
  
Great, junior 2240 did that and told me '10'  
That little punk gave me a non-zero return.  
I'm glad he's dead  
>
```


Answers: Child Exit Status

```
1 // child_wait_returnval.c: fork/exec plus parent waits for child and
2 // checks their status using macros. If nonzero, parent reports.
3
4 #include <stdio.h>
5 #include <stdlib.h>
6 #include <sys/wait.h>
7 #include <unistd.h>
8
9 int main(int argc, char* argv){
10     char *child_argv[] = {"/complain",NULL};           // program returns non-zero
11     char *child_cmd = "complain";
12
13     printf("I'm %d, and I really don't feel like '%s'ing\n",
14           getpid(),child_cmd);
15     printf("I have a solution\n");
16
17     pid_t child_pid = fork();
18
19     if(child_pid == 0){
20         printf(" I'm %d My pa %d' wants me to '%s'. This sucks.\n",
21               getpid(), getppid(), child_cmd);
22         execvp(child_cmd, child_argv);
23         printf(" I don't feel like myself anymore...\n"); // unreachable
24     }
25     else{
26         int status;
27         wait(&status);                                // wait for child to finish, collect status
28         if(WIFEXITED(status)){
29             int retval = WEXITSTATUS(status);          // decode status to 0-255
30             printf("Great, junior %d did that and told me %d'\n",
31                   child_pid, retval);
32             if(retval != 0){                            // nonzero exit codes usually indicate failure
33                 printf("That little punk gave me a non-zero return. I'm glad he's dead\n");
34             }
35         }
36     }
37     return 0;
38 }
```

Return Value for wait() family

- ▶ Return value for wait() and waitpid() is the PID of the child that finished
- ▶ Makes a lot of sense for wait() as multiple children can be started and wait() reports which finished
- ▶ One wait() per child process is typical
- ▶ See faster_child.c

```
// parent waits for each child
for(int i=0; i<3; i++){
    int status;
    int child_pid = wait(&status);
    if(WIFEXITED(status)){
        int retval = WEXITSTATUS(status);
        printf("PARENT: Finished child proc %d, retval: %d\n",
               child_pid, retval);
    }
}
```

Blocking vs. Nonblocking Activities

Blocking

- ▶ A call to `wait()` and `waitpid()` may cause calling process to **block** (hang, stall, pause, suspend, so many names...)
- ▶ Blocking is associated with other activities as well
 - ▶ I/O, obtain a lock, get a signal, etc.
- ▶ Generally creates **synchronous** situations: waiting for something to finish means the next action *always* happens..
next

```
// BLOCKING VERSION
```

```
int pid = waitpid(child_pid, &status, 0);
```

Non-blocking

- ▶ Contrast with **non-blocking** (asynchronous) activities: calling process goes ahead even if something isn't finished yet
- ▶ `wait()` is always blocking
- ▶ `waitpid()` can be blocking or non-blocking

Non-Blocking waitpid()

- ▶ Use the WNOHANG option
- ▶ Returns immediately regardless of the child's status

```
int child_pid = fork();
int status;

// NON-BLOCKING
int pid = waitpid(child_pid, &status, WNOHANG); // specific child
OR
int pid = waitpid(-1, &status, WNOHANG); // any child
```

Returned pid is

Returned	Means
child_pid	status of child has changed (exit)
0	there is no status change for child
-1	an error

Examine `impatient_parent.c`

impatient_parent.c


```
1 // impatient_parent.c: demonstrate non-blocking waitpid(),
2
3 #include <stdio.h>
4 #include <stdlib.h>
5 #include <sys/wait.h>
6 #include <unistd.h>
7
8 int main(int argc, char* argv){
9
10     char *child_argv[] = {"/complain",NULL};
11     char *child_cmd = "complain";
12
13     printf("PARENT: Junior is about to '%s', I'll keep an eye on him\n",
14           child_cmd);
15
16     pid_t child_pid = fork();
17
18     // CHILD CODE
19     if(child_pid == 0){
20         printf("CHILD: I'm %d and I'm about to '%s'\n",
21               getpid(), child_cmd);
22         execvp(child_cmd, child_argv);
23     }
24
25     // PARENT CODE
26     int status;
27     int pid = waitpid(child_pid,&status,WNOHANG); // Check if child done, but don't actually wait
28     if(pid == child_pid && WIFEXITED(status)){ // Child did finish
29         printf("PARENT: Good job junior. You told me %d\n",WEXITSTATUS(status));
30     }
31     else{ // Child not done yet
32         printf("PARENT: %d? The kid's not done yet. I'm bored\n",pid);
33     }
34     return 0;
35 }
```

Runs of impatient_parent.c

```
> gcc impatient_parent.c
> a.out
PARENT: Junior is about to 'complain', I'll keep an eye on him
PARENT: 0? The kid's not done yet. I'm bored
CHILD: I'm 1863 and I'm about to 'complain'
> COMPLAIN: God this sucks. On a scale of 0 to 10 I hate pa ...

> a.out
PARENT: Junior is about to 'complain', I'll keep an eye on him
PARENT: 0? The kid's not done yet. I'm bored
CHILD: I'm 1865 and I'm about to 'complain'
> COMPLAIN: God this sucks. On a scale of 0 to 10 I hate pa ...
```

Exercise: Helicopter Parent



- ▶ Modify `impatient_parent.c` to `helicopter_parent.c`
- ▶ Checks continuously on child process
- ▶ Will need a loop for this...

```
> gcc helicopter_parent.c
> a.out
PARENT: Junior is about to 'complain', I'll keep an eye on him
Oh, junior's taking so long. Is he among the 50% of people that are below average?
Oh, junior's taking so long. Is he among the 50% of people that are below average?
...
Oh, junior's taking so long. Is he among the 50% of people that are below average?
Oh, junior's taking so long. Is he among the 50% of people that are below average?
CHILD: I'm 21789 and I'm about to 'complain'
Oh, junior's taking so long. Is he among the 50% of people that are below average?
...
Oh, junior's taking so long. Is he among the 50% of people that are below average?
Oh, junior's taking so long. Is he among the 50% of people that are below average?
COMPLAIN: God this sucks. On a scale of 0 to 10 I hate pa ...
Oh, junior's taking so long. Is he among the 50% of people that are below average?
Oh, junior's taking so long. Is he among the 50% of people that are below average?
...
PARENT: Good job junior. I only checked on you 226 times.
```

Answers: Helicopter Parent

```
1 // demonstrate non-blocking waitpid() in excess
2 #include <stdio.h>
3 #include <stdlib.h>
4 #include <sys/wait.h>
5 #include <unistd.h>
6
7 int main(int argc, char* argv){
8
9     char *child_argv[] = {"/complain",NULL};
10    char *child_cmd = "complain";
11
12    printf("PARENT: Junior is about to '%s', I'll keep an eye on him\n",
13           child_cmd);
14
15    pid_t child_pid = fork();
16
17    // CHILD CODE
18    if(child_pid == 0){
19        printf("CHILD: I'm %d and I'm about to '%s'\n",
20               getpid(), child_cmd);
21        execvp(child_cmd, child_argv);
22    }
23
24    // PARENT CODE
25    int status;
26    int checked = 0;
27    while(1){
28        int cpid = waitpid(child_pid,&status,WNOHANG); // Check if child done, but don't actually wait
29        if(cpid == child_pid){ // Child did finish
30            break;
31        }
32        printf("Oh, junior's taking so long. Is he among the 50%% of people that are below average?\n");
33        checked++;
34    }
35    printf("PARENT: Good job junior. I only checked on you %d times.\n",checked);
36    return 0;
37 }
```


Polling vs Interrupts

- ▶ `helicopter_parent.c` is an example of **polling**: checking on something repeatedly until it achieves a ready state
- ▶ Easy to program, generally inefficient
- ▶ Alternative: **interrupt** style is closer to `wait()` and `waitpid()` *without* `WNOHANG`: rest until notified of a change
- ▶ Usually requires cooperation with OS/hardware which must wake up process when stuff is ready
- ▶ Both polling-style and interrupt-style programming have uses

Zombies...

- ▶ Parent creates a child
- ▶ Child completes
- ▶ Child becomes a **zombie** (!!!)
- ▶ Parent waits for child
- ▶ Child eliminated



Didn't see that coming next, did you?

Zombie Process

A process that has finished, but has not been `wait()`'ed for by its parent yet so cannot be eliminated from the system. OS can reclaim child resources like memory once parent `wait()`'s.

Demonstrate

Requires a process monitoring with `top/ps` but can see zombies created using `spawn_undead.c`

Tree of Processes

```

    pstree
systemd+-NetworkManager---2*[{NetworkManager}]
|accounts-daemon---2*[{accounts-daemon}]
|colord---2*[{colord}]
|csd-printer---2*[{csd-printer}]
|cupsd
|dbus-daemon
|drjava---java+-java---27*[{java}]
|      ^-37*[{java}]
|dropbox---106*[{dropbox}]
|emacs+-aspell
|      |bash---pstree
|      |evince---4*[{evince}]
|      |idn
|      ^-3*[{emacs}]
|gdm+-gdm-session-wor+-gdm-wayland-ses+-gnome-session-b+-gnome-shell+-Xwayland---14*[{Xwayland}]
|...
|      |gnome-terminal+-bash+-chromium+-chrome-sandbox---chromium---chromium+-8*[{chromium---12*[{chromium}]]]
|      |      |chromium---11*[{chromium}]
|      |      |chromium---14*[{chromium}]
|      |      |chromium---15*[{chromium}]
|      |      ^-chromium---18*[{chromium}]
|      |      |chromium---9*[{chromium}]
|      |      ^-42*[{chromium}]
|      |      ^-cinnamon---21*[{cinnamon}]
|      |      |bash---ssh
|      |      ^-3*[{gnome-terminal-}]

```

- ▶ Processes exist in a tree: see with shell command `ps tree`
- ▶ Children can be **orphaned** by parents: parent exits without `wait()`'ing for child
- ▶ Orphans are adopted by the root process
 - ▶ `init` traditionally
 - ▶ `systemd` in many modern systems
- ▶ Root process occasionally `wait()`'s to "reap" zombies