

# CSCI 4061: Files, Directories, Standard I/O

Chris Kauffman

*Last Updated:  
Tue Feb 19 09:34:42 CST 2019*

# Logistics

## Reading

- ▶ Stevens/Rago  
Ch 3, 4, 5, 6
- ▶ OR Robbins and Robbins  
Ch 4, 5

## Goals

- ▶ Std I/O vs Unix Syscall
- ▶ File / Directory Functions
- ▶ Filesystem

## Lab04: Pipes

How did it go?

## Project 1

Questions?

## Exam 1: Next week

- ▶ Tue Review
- ▶ Thu Exam

## Exercise: Quick Recap

1. What is a pipe? What system call is used to create it?  
Example?
2. How does one put data into a pipe? Get data from a pipe?
3. How can one arrange for communication between a parent and child process?
  - ▶ Child to parent
  - ▶ Parent to child
4. What syntax do standard shells use to redirect program output to files?
5. What low-level system calls are used to accomplish redirection?

# Answers: Quick Recap

1. What is a pipe? What system call is used to create it? Example?
  - ▶ Internal OS communication buffer, created via

```
int pip;  
int result = pipe(pip);
```
2. How does one put data into a pipe? Get data from a pipe?
  - ▶ 

```
nbytes = write(pip[1], w_buff, BUFLLEN);  
nbytes = read(pip[0], r_buff, BUFLLEN);
```
3. How can one arrange for communication between a parent and child process?
  - ▶ Child to parent: parent opens pipe, child writes, parent reads
  - ▶ Parent to child: parent opens pipe, parent writes, child reads
4. What syntax do standard shells use to redirect program output to files? Read input from files?
  - ▶ 

```
$> my_program arg1 arg2 > output.txt  
$> other_prog arg1 < input.txt
```
5. What low-level system calls are used to accomplish redirection?
  - ▶ 

```
dup2(fd_a, fd_b);
```
  - ▶ writes to `fd_b` write to `fd_a` instead
  - ▶ reads from `fd_b` read from `fd_a` instead

# Permissions / Modes

- ▶ Unix enforces file security via *modes*: permissions as to who can read / write / execute each file
- ▶ See permissions/modes with `ls -l`
- ▶ Look for series of 9 permissions

```
> ls -l
total 140K
-rwx--x--- 2 kauffman faculty 8.6K Oct  2 17:39 a.out
-rw-r--r-- 1 kauffman devel  1.1K Sep 28 13:52 files.txt
-rw-rw---- 1 kauffman faculty 1.5K Sep 26 10:58 gettysburg.txt
-rwx--x--- 2 kauffman faculty 8.6K Oct  2 17:39 my_exec
----- 1 kauffman kauffman 128 Oct  2 17:39 unreadable.txt
-rw-rw-r-x 1 root      root      1.2K Sep 26 12:21 scripty.sh
  U  G  O      O      G      S      M  T      N
  S  R  T      W      R      I      O  I      A
  E  O  H      N      O      Z      D  M      M
  R  U  E      E      U      E      E      E
      P  R      R      P
```

~~~~~  
PERMISSIONS

- ▶ Every file has permissions set from somewhere on creation

## Changing Permissions

Owner of file (and sometimes group member) can change permissions via `chmod`

```
> ls -l a.out
```

```
-rwx--x--- 2 kauffman faculty 8.6K Oct 2 17:39 a.out
```

```
> chmod u-w,g+r,o+x a.out
```

```
> ls -l a.out
```

```
-r-xr-x--x 2 kauffman faculty 8.6K Oct 2 17:39 a.out
```

- ▶ `chmod` also works via octal bits (suggest against this unless you want to impress folks at parties)
- ▶ Programs specify permissions for files they create via C calls
- ▶ Curtailed by the `umask` shell or `umask()` C function: indicates permissions that are not allowed
- ▶ Common program strategy: create files with very liberal read/write/execute permissions, `umask` of user will limit this

## Exercise: Regular File Creation Basics

### C Standard I/O

- ▶ Write/Read data?
- ▶ Open a file, create it if needed?
- ▶ Result of opening a file?
- ▶ Close a file?
- ▶ Set permissions on file creation?

### Unix System Calls

- ▶ Write/Read data?
- ▶ Open a file, create it if needed?
- ▶ Result of opening a file?
- ▶ Close a file?
- ▶ Set permissions on file creation?

# Answers: Regular File Creation Basics

## C Standard I/O

- ▶ Write/Read data?

```
fscanf(), fprintf()  
fread(), fwrite()
```

- ▶ Open a file, create it if needed?
- ▶ Result of opening a file?

```
FILE *out =  
    fopen("myfile.txt", "w");
```

- ▶ Close a file?  
  
`fclose(out);`
- ▶ Set permissions on file creation?  
Not possible... dictated by `umask`

## Unix System Calls

- ▶ Write/Read data?

```
write(), read()
```

- ▶ Open a file, create it if needed?
- ▶ Result of opening a file?

```
int fd =  
    open("myfile.txt",  
        O_WRONLY | O_CREAT,  
        permissions);
```

- ▶ Close a file?  
  
`close(fd);`
- ▶ Set permissions on file creation?
  - ▶ Additional options to `open()`, which brings us to...



## Permissions / Modes in C Calls

- ▶ Default `open(name,opts)` has NO PERMISSIONS
- ▶ When opening with `O_CREAT`, specify permissions for new file
- ▶ `int fd = open(name, opts, mode);`

| Symbol  | Entity | Sets    |
|---------|--------|---------|
| S_IRUSR | User   | Read    |
| S_IWUSR | User   | Write   |
| S_IXUSR | User   | Execute |
| S_IRGRP | Group  | Read    |
| S_IWGRP | Group  | Write   |
| S_IXGRP | Group  | Execute |
| S_IROTH | Others | Read    |
| S_IWOTH | Others | Write   |
| S_IXOTH | Others | Execute |

### Compare: `write_readable.c` VERSUS `write_unreadable.c`

```
char *outfile = "newfile.txt";           // doesn't exist yet
int flags      = O_WRONLY | O_CREAT;      // write/create
mode_t perms   = S_IRUSR | S_IWUSR;       // variable for permissions
int out_fd     = open(outfile, flags, perms);
                ~~~~~
```

# C Standard I/O Implementation

Typical Unix implementation of standard I/O library FILE is

- ▶ A file descriptor
- ▶ Some buffers with positions
- ▶ Some options controlling buffering

From `/usr/lib/libio.h`

```
struct _IO_FILE {  
    int _flags;                // options  
    char* _IO_read_ptr;        // positions and  
    char* _IO_read_end;        // buffers for  
    char* _IO_read_base;        // read and write  
    char* _IO_write_base;  
    ...;  
    int _fileno;                // file descriptor  
    ...;  
    _IO_lock_t *_lock;          // locking  
};
```

## Exercise: Subtleties of Mixing Standard and Low-Level I/O

- ▶ Predict output of program given input file
- ▶ Use knowledge that buffering occurs internally for standard I/O library
- ▶ **Note:** Similar subtleties exist if FILE\* are not properly closed
- ▶ FILE buffers may contain unflushed data: not written at close
- ▶ See fail-to-write.c
- ▶ File descriptors always get flushed out by OS

3K.txt:

```
1 2 3 4 5 6 7 8 9 10 11 12 13 14...
37 38 39 40 41 42 43 44 45 46 47 ...
70 71 72 73 74 75 76 77 78 79 80 ...
102 103 104 105 106 107 108 109 1...
...
```

mixed-std-low.c:

```
1 int main(int argc, char *argv[]){
2
3     FILE *input = fopen("3K.txt","r");
4     int first;
5     fscanf(input, "%d", &first);
6     printf("FIRST: %d\n",first);
7
8     int fd = fileno(input);
9     char *buf[64];
10    read(fd, buf, 63);
11    buf[127] = '\0';
12    printf("NEXT: %s\n",buf);
13
14    return 0;
15 }
```

# Controlling FILE Buffering

```
#include <stdio.h>
void setbuf(FILE *stream, char *buf);
void setbuffer(FILE *stream, char *buf, size_t size);
void setlinebuf(FILE *stream);
int setvbuf(FILE *stream, char *buf, int mode, size_t size);
```

Series of functions which control buffering. Example:

```
// Turn off buffering of stdout
setvbuf(stdout, NULL, _IONBF, 0);
```

Why should this line be familiar to **ALL** of you?

# Filesystems, inodes, links

- ▶ Unix **filesystems** implement physical layout of files/directories on a storage media (disks, CDs, etc.)
- ▶ Many filesystems exist but all Unix-centric filesystems share some common features

## inode

- ▶ Data structure which describes a single file
- ▶ Stores some meta data: inode#, size, timestamps, owner
- ▶ A table of contents: which disk blocks contain file data
- ▶ Does **not** store filename, does store a **link count**

## Directories

- ▶ List names and associated inode
- ▶ Each entry constitutes a **hard link** to an inode or a **symbolic link** to another file
- ▶ Files with 0 hard links are deleted

# Rough Filesystem in Pictures

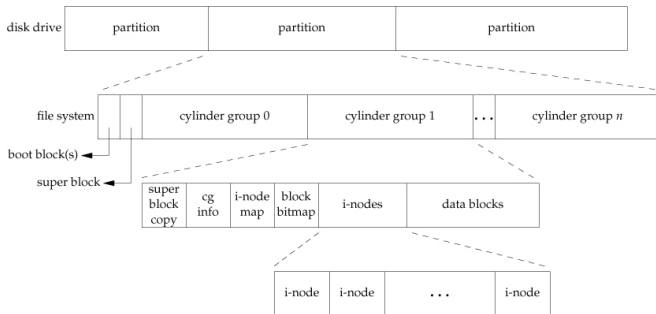
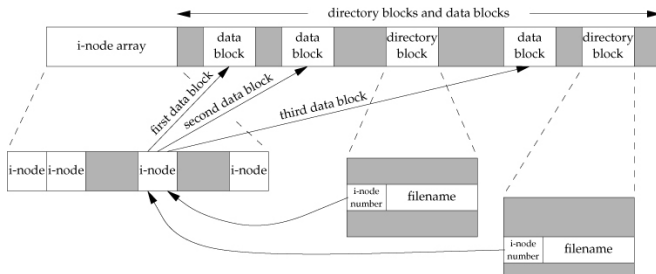


Figure 4.13 Disk drive, partitions, and a file system (Stevens/Rago)



# Shell Demo of Hard and Symbolic Links

```
> rm *
> touch fileX          # create empty fileX
> touch fileY          # create empty fileY
> ln fileX fileZ        # hard link to fileX called fileZ
> ln -s fileX fileW     # symbolic link to fileX called fileW
> ls -li               # -i for inode numbers
total 12K
6685588 -rw-rw---- 2 kauffman kauffman 0 Oct  2 21:24 fileX
6685589 -rw-rw---- 1 kauffman kauffman 0 Oct  2 21:24 fileY
6685588 -rw-rw---- 2 kauffman kauffman 0 Oct  2 21:24 fileZ
6685591 lrwxrwxrwx 1 kauffman kauffman 5 Oct  2 21:29 fileB -> fileA
6685590 lrwxrwxrwx 1 kauffman kauffman 5 Oct  2 21:25 fileW -> fileX
↑↑↑↑↑↑↑ ↑          ↑          ↑↑↑↑↑↑↑↑
inode#  regular    hard link count          symlink target
        or symlink

> file fileW           # file type of fileW
fileW: symbolic link to fileX
> file fileB           # file type of fileB
fileB: broken symbolic link to fileA
```

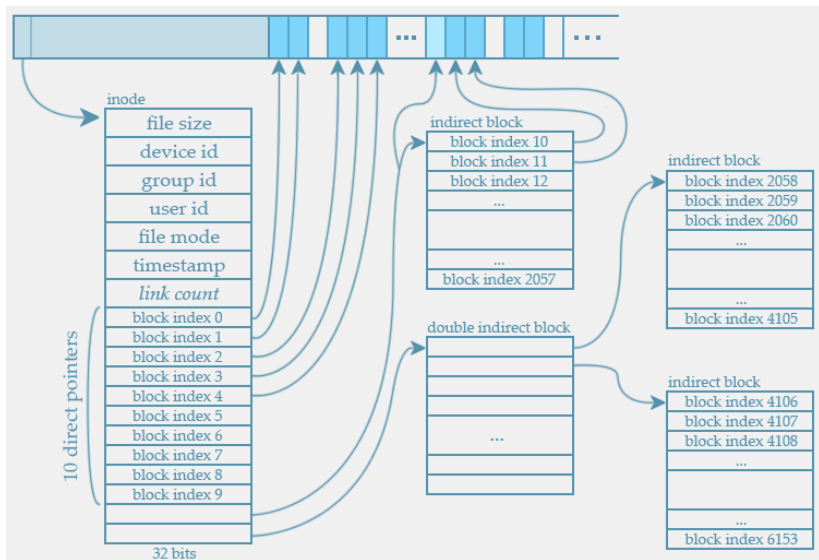
# Linking Commands and Functions

| Shell Command                  | C Function                                                     | Effect                                                          |
|--------------------------------|----------------------------------------------------------------|-----------------------------------------------------------------|
| <code>ln fileX fileY</code>    | <code>link("fileX", "fileY");</code>                           | Create a hard link                                              |
| <code>rm fileX</code>          | <code>remove("fileX");</code><br><code>unlink("fileX");</code> | Unlink (remove) hard link<br>Identical to <code>remove()</code> |
| <code>ln -s fileX fileY</code> | <code>symlink("fileX", "fileY");</code>                        | Create a Symbolic link                                          |

- ▶ Creating hard links preserves inodes
- ▶ Hard links not allowed for directories unless you are root
  - > `ln /home/kauffman to-home`
  - `ln: /home/kauffman: hard link not allowed for directory`
  - Can create directory cycles if this was allowed
- ▶ Symlinks easily identified so utilities can skip them



# FYI: inodes are a complex beast themselves



Source: File System Design by Justin Morgan

## sync() and Internal OS Buffers

- ▶ Operating system maintains internal data associated with open files
- ▶ Writing to a file doesn't go immediately to a disk
- ▶ May live in an internal buffer for a while before being sync'ed to physical medium (OS buffer cache)

| Shell Command | C function  | Effect                                                    |
|---------------|-------------|-----------------------------------------------------------|
| sync          | sync();     | Synchronize cached writes to persistent storage           |
|               | syncfs(fd); | Synchronize cached writes for filesystem of given open fd |

- ▶ Sync called so that one can "Safely remove drive"
- ▶ Sync happens automatically at regular intervals (ex: 15s)

# Basic File Statistics via stat

| Command   | C function                     | Effect                  |
|-----------|--------------------------------|-------------------------|
| stat file | int ret = stat(file,&statbuf); | Get statistics on file  |
|           | int fd = open(file,...);       | Same as above but with  |
|           | int ret = fstat(fd,&statbuf);  | an open file descriptor |

Shell command stat provides basic file info such as shown below

```
> stat a.out
  File: a.out
  Size: 12944      ^IBlocks: 40          IO Block: 4096   regular file
Device: 804h/2052d^IInode: 6685354      Links: 1
Access: (0770/-rwxrwx---)  Uid: ( 1000/kauffman)   Gid: ( 1000/kauffman)
Access: 2017-10-02 23:03:21.192775090 -0500
Modify: 2017-10-02 23:03:21.182775091 -0500
Change: 2017-10-02 23:03:21.186108423 -0500
Birth: -

> stat /
  File: /
  Size: 4096      ^IBlocks: 8          IO Block: 4096   directory
Device: 803h/2051d^IInode: 2          Links: 17
Access: (0755/drwxr-xr-x)  Uid: (   0/   root)   Gid: (   0/   root)
Access: 2017-10-02 00:56:47.036241675 -0500
Modify: 2017-05-07 11:34:37.765751551 -0500
Change: 2017-05-07 11:34:37.765751551 -0500
Birth: -
```

See stat-demo.c for info on C calls to obtain this info

# Directory Access

- ▶ Directories are fundamental to Unix (and most file systems)
- ▶ Unix file system rooted at / (root directory)
- ▶ Subdirectories like bin, ~/home, and /home/kauffman
- ▶ Useful shell commands and C function calls pertaining to directories are as follows

| Shell Command | C function                          | Effect                           |
|---------------|-------------------------------------|----------------------------------|
| mkdir name    | int ret = mkdir(path,perms);        | Create a directory               |
| rmdir name    | int ret = rmdir(path);              | Remove empty directory           |
| cd path       | int ret = chdir(path);              | Change working directory         |
| pwd           | char *path = getcwd(buf,SIZE);      | Current directory                |
| ls            | DIR *dir = opendir(path);           | List directory contents          |
|               | struct dirent *file = readdir(dir); | Start reading filenames from dir |
|               | int ret = closedir(dir);            | Call in a loop, NULL when done   |
|               |                                     | After readdir() returns NULL     |

See `dir-demo.c` for demonstrations

## Movement within Files

- ▶ Can move OS internal position in a file around with `lseek()`
- ▶ Note that size is arbitrary: can seek to any positive position
- ▶ File automatically expands if position is larger than current size - fills holes with 0s (null chars)
- ▶ Examine `file-hole.c` and `file-hole2.c`

| C function                                        | Effect                   |
|---------------------------------------------------|--------------------------|
| <code>int res = lseek(fd, offset, option);</code> | Move position in file    |
| <code>lseek(fd, 20, SEEK_CUR);</code>             | Move 20 bytes forward    |
| <code>lseek(fd, 50, SEEK_SET);</code>             | Move to position 50      |
| <code>lseek(fd, -10, SEEK_END);</code>            | Move 10 bytes from end   |
| <code>lseek(fd, +15, SEEK_END);</code>            | Move 15 bytes beyond end |

See also C standard I/O `fseek(FILE *)` / `rewind(FILE *)` functions

## fcntl(): Jack of all trades

- ▶ fcntl() does a bunch of stuff
- ▶ Some previous calls implemented with fcntl()
  - ▶ `int fd2 = dup(fd1);` OR
  - ▶ `int fd2 = fcntl(fd1, F_DUPFD);`

```
#include <fcntl.h>
#include <unistd.h>
#include <sys/types.h>
```

```
int fcntl(int fd, int cmd, /* arg */ ...);
```

| Command        | Effect                                                          |
|----------------|-----------------------------------------------------------------|
| F_DUPFD        | duplicate a file descriptor                                     |
| F_GETFD        | get file descriptor flags                                       |
| F_SETFD        | set file descriptor flags                                       |
| F_GETFL        | get file status flags and access modes                          |
| F_SETFL        | set file status flags and access modes                          |
| F_GETOWN       | get proc ID currently receiving SIGIO and SIGURG signals for fd |
| F_SETOWN       | set proc ID that will receive SIGIO and SIGURG signals for fd   |
| <b>Locking</b> |                                                                 |
| F_GETLK        | get first lock that blocks description specified by arg         |
| F_SETLK        | set or clear segment lock specified by arg                      |
| F_SETLKW       | same as FSETLK except it blocks until request satisfied         |
| ...            |                                                                 |

## `select()` and `poll()`: Non-busy waiting

- ▶ Recall **polling** is a busy wait on something: constantly check until ready
- ▶ Alternative is **interrupt-driven** wait: ask for notification when something is ready, go to sleep, get woken up
- ▶ Waiting is often associated with input from other processes through pipes or sockets
- ▶ Both `select()` and `poll()` allow for waiting on input from multiple file descriptors
- ▶ Confusingly, **both `select()` and `poll()` are interrupt-driven**: will put process to sleep until something changes in one or more files
- ▶ `poll()` doesn't do polling (busy wait) - it does interrupt driven I/O (!!)
- ▶ Example application: database system is waiting for any of 10 users to enter a query, don't know which one will type first

# File Descriptor Sets

- ▶ `select()` uses file descriptor **sets**
- ▶ `fd_set` tracks descriptors of interest, operated on with macros

```
fd_set my_set;  
void FD_ZERO(fd_set *set);           // clear entire set  
void FD_SET(int fd, fd_set *set);    // fd now in set  
void FD_CLR(int fd, fd_set *set);    // fd now not in set  
int  FD_ISSET(int fd, fd_set *set);  // test if fd in set
```

- ▶ Example: setup set of potential read sources

```
int pipeA[2], pipeB[2], rd_fd;      // set up several read sources  
pipe(pipeA);  
pipe(pipeB);  
rd_fd = open("myfile.txt",RD_ONLY);  
  
fd_set read_set;                    // set of file descriptors for select()  
FD_ZERO(&read_set);                // init the set  
  
FD_SET(pipeA[PREAD], &read_set);    // include read ends of pipes in set  
FD_SET(pipeB[PREAD], &read_set);  
FD_SET(rd_fd, &read_set);           // include read file in the set
```



## Multiplexing: Efficient input from multiple sources

- ▶ `select()` block a process until at least one of member of the `fd_set` is "ready"
- ▶ Most common use: waiting for input from multiple sources
- ▶ Example: Multiple child processes writing to pipes at different rates

```
#include <sys/select.h>
fd_set read_set, write_set,      // sets of fds to wake up for
    except_set;

struct timeval timeout;          // allows timeout: wake up if nothing happens

int nfds =                       // returns nfds changed
    select(maxfd+1,              // must pass max fd+1
        &read_set,              // any of set may be NULL to ignore
        &write_set,
        &except_set,
        &timeout);              // NULL time waits indefinitely
```

- ▶ Lab07 covers `select()` with two children
- ▶ See `select-pipes.c` shows multiple children with different communication rates