

CSCI 4061: Signals and Signal Handlers

Chris Kauffman

*Last Updated:
Tue Mar 12 05:16:00 CDT 2019*

Logistics

Reading

- ▶ Stevens/Rago
Ch 10
- ▶ OR Robbins and Robbins
Ch 8.1-8.7, 9.1-2

Goals

- ▶ Sending Signals in C
- ▶ Signal Handlers
- ▶ `select()`: Multiplexing I/O

Lab07: pmap / signals intro

How did it go?

Project 2

- ▶ Under development
- ▶ Will discuss on Tue

Exercise: Lab07 Signals

1. What is a signal?
2. What system call is used to send a process a signal? How is it invoked?

Answers: Lab07 Signals

1. What is a signal?
 - ▶ Notification from somewhere, limited information, special effects
2. What system call is used to send a process a signal? How is it invoked?
 - ▶ `kill(pid, SIGSOMTHING);`

What kind of signals are there?

- ▶ Signals are an old system of communication to convey a limited amount of info to a process
- ▶ "Delivered" by the OS to a running process to inform of it of an event
- ▶ Process responds in one of several ways according to its **disposition**
- ▶ **Asynchronous**: could delivered to a process **at any time**

Process Signal Disposition

```
> man 7 signal
```

```
...
```

Signal dispositions

Each signal has a current disposition, which determines how the process behaves when it is delivered the signal.

The entries in the "Action" column of the tables below specify the default disposition for each signal, as follows:

Term Default action is to terminate the process.

Ign Default action is to ignore the signal.

Core Default action is to terminate the process and dump core (see core(5)).

Stop Default action is to stop the process.

Cont Default action is to continue the process if it is currently stopped.

Can be adjust signal disposition with various system calls to establish **signal handlers** for the process.

Standard Types of Signals

```
> man 7 signal
Standard Signals
```

Signal	x86 Value	Default Action	Comment
SIGHUP	1	Term	Hangup detected on controlling terminal or death of controlling process
SIGINT	2	Term	Interrupt from keyboard
SIGQUIT	3	Core	Quit from keyboard
SIGILL	4	Core	Illegal Instruction
SIGTRAP	5	Core	Trace/breakpoint trap
SIGABRT	6	Core	Abort signal from abort(3)
SIGBUS	7	Core	Bus error (bad memory access)
SIGFPE	8	Core	Floating-point exception (CK: actually integer divide by 0)
SIGKILL	9	Term	Kill signal
SIGUSR1	10	Term	User-defined signal 1
SIGSEGV	11	Core	Invalid memory reference
SIGUSR2	12	Term	User-defined signal 2
SIGPIPE	13	Term	Broken pipe: write to pipe with no readers; see pipe(7)
SIGALRM	14	Term	Timer signal from alarm(2)
SIGTERM	15	Term	Termination signal
SIGSTKFLT	16	Term	Stack fault on coprocessor (unused)
SIGCHLD	17	Ign	Child stopped or terminated
SIGCONT	18	Cont	Continue if stopped
SIGSTOP	19	Stop	Stop process
SIGTSTP	20	Stop	Stop typed at terminal
...			
SIGUNUSED	31	Core	Synonymous with SIGSYS

Note: Different CPU architectures may have different values for some signals and support other signals not listed

(Ex: MIPS CPUs use SIGCONT=25 with a synonym for SIGCHLD=19)

Basic Signal Handlers via `signal()`

Pressing Ctrl-c in a terminal sends SIGINT to a running program which normally Terminates the program. The below template establishes a **signal handler** for SIGINT.

```
#include <signal.h>
void handle_SIGINT(int sig_num) {
    ...
}

int main () {
    // Set handling functions for programs
    signal(SIGINT, handle_SIGINT);
    ...
}
```

- ▶ When SIGINT arrives at program, control jumps to function `handle_SIGINT()` with argument `sig_num == SIGINT`
- ▶ When `handle_SIGINT()` completes, control returns to wherever the program left off

Examine: `no_interruptions_signal.c`

History Note: Resetting Signal Handlers

```
void handle_SIGINT(int sig_num) {  
    signal(SIGINT, handle_SIGINT);  
    // Reset handler to catch SIGINT next time  
    // Not needed in modern systems  
    printf("\nNo SIGINT-erruptions allowed.\n");  
    fflush(stdout);  
}  
  
int main () {  
    signal(SIGINT, handle_SIGINT);  
    ...  
}
```

- ▶ Old sources describe the need to reset handles while running
- ▶ Why is this subtly awful?
- ▶ Not needed on *most* modern Unix systems

Historical Notes

- ▶ Signals were an early concept but were initially "unreliable": might get lost and so were not as useful as their modern incarnation
- ▶ Historically, required to reset signal handlers after they were called. First line of handler was always `signal(this_signal, this_handler);` though this was still buggy.
- ▶ Historically, some system calls could be interrupted by signals. Robbins & Robbins go on and on about this.

On FreeBSD 8.0, Linux 3.2.0, and Mac OS X 10.6.8, when signal handlers are installed with the `signal` function, interrupted system calls will be restarted. The default on Solaris 10, however, is to return an error (`EINTR`) instead when system calls are interrupted by signal handlers installed with the `signal` function.

– Stevens and Rago, 10.5

Portability Notes on `signal()`

```
> man 2 signal
```

```
...
```

The behavior of `signal()` varies across UNIX versions, and has also varied historically across different versions of Linux.

AVOID ITS USE: use `sigaction(2)` instead.

PORTABILITY

The semantics when using `signal()` to establish a signal handler vary across systems (and POSIX.1 explicitly permits this variation); **do not use it for this purpose.**

- ▶ `signal()` part of the C standard but is old with different behaviors across different systems
- ▶ POSIX defined new functions which were designed to break from its tradition and fix problems associated with it
- ▶ Requires introduction of **signal sets**, data type for a set of signals along with associated functions

Portable Signal Handlers via sigaction()

- ▶ The sigaction() function is more portable than signal() to register signal handlers.
- ▶ Makes use of struct sigaction which specifies properties of signal handler registrations, most importantly the field **sa_handler**

```
int main(){                                     // SAMPLE HANDLER SETUP USING sigaction()
    struct sigaction my_sa = {};               // portable signal handling setup with sigaction()
    my_sa.sa_handler = handle_signals;         // run function handle_signals
    sigemptyset(&my_sa.sa_mask);              // don't block any other signals during handling
    my_sa.sa_flags = SA_RESTART;               // restart system calls on signals if possible
    sigaction(SIGTERM, &my_sa, NULL);         // register SIGTERM with given action
    sigaction(SIGINT, &my_sa, NULL);          // register SIGINT with given action
    ...;
}
```

See `no_interruptions_sigaction.c`

Ignoring Signals, Restoring Defaults

- ▶ Setting the signal handler to `SIG_IGN` will cause signals to be silently ignored.
- ▶ Setting the signal handler to `SIG_DFL` will restore default disposition.

Demo `no_interruptions_ignore.c`

Sleeping, Pausing, and Stopping

Sleeping/Pausing: wait for a signal

- ▶ `sleep(5)` suspends process execution until a signal is delivered or for 5 seconds elapses
- ▶ `pause()` suspends process execution until a signal is delivered;
- ▶ `sleep(0)` is equivalent to `pause()`

Note sleep behavior of various `no_interruptions` programs

Signals that Affect Execution

- ▶ `SIGSTOP` will causes process to stop, will not resume until...
- ▶ `SIGCONT` causes a stopped process to resume, otherwise ignored by default
- ▶ All signals are delivered while a process is stopped BUT it is not resumed until receiving `SIGCONT`

Examine: `start_stop.c` with `circle_of_life.c`

You want the Signal? You Can't Handle the Signal!

- ▶ SIGKILL and SIGSTOP cannot be handled differently from default
 - ▶ SIGKILL always terminates a process
 - ▶ SIGSTOP always stops a process execution
- ▶ In that sense they are a little different than the other signals but use the same OS delivery mechanism and `kill()` semantics
- ▶ Calls to `sigaction()` or `signal()` for these two will fail
- ▶ See `cant_handle_kill.c`

Other Parts of struct sigaction

Type	Field	Purpose
<code>void(*) (int)</code>	<code>sa_handler</code>	Pointer to a signal-catching function or one of the macros <code>SIG_IGN</code> or <code>SIG_DFL</code> .
<code>sigset_t</code>	<code>sa_mask</code>	Additional set of signals to be blocked during execution of signal-catching function.
<code>int</code>	<code>sa_flags</code>	Special flags to affect behavior of signal. Typically <code>SA_RESTART</code> is used to restart system calls automatically
<code>...</code>	<code>sa_sigaction</code>	More complex handler used when <code>sa_flags</code> has <code>SA_SIGINFO</code> set; passes additional info to handler like PID of signaling process.

Standard setup for `sigaction()` call is

```
struct sigaction my_sa = {};  
sigemptyset(&my_sa.sa_mask);           // don't block any other signals during handl  
my_sa.sa_flags = SA_RESTART;            // always restart system calls on signals poss  
my_sa.sa_handler = handle_SIGTERM;      // run function handle_SIGTERM  
sigaction(SIGTERM, &my_sa, NULL);       // register SIGTERM with given action
```

Dangers in Signal Handlers

- ▶ General advice: do as little as possible in a signal handler
- ▶ Make use of only **reentrant** functions
 - ... reentrant if it can be interrupted in the middle of its execution, and then be safely called again ("re-entered") before its previous invocations complete execution.*
 - [Wikipedia: Reentrancy](#)
- ▶ Notably not reentrant
 - `printf()` family, `malloc()`, `free()`
- ▶ Reentrant functions pertinent to thread-based programming as well (later)
- ▶ Demo `non-reentrant.c`

Exercise: Non-Reentrant Function Example

- ▶ Program calls non-reentrant function `f()` in `main()` and `handle_signal()`
- ▶ With no interrupts, would expect to see 7 printed, with interrupts see 19,7 in either order
- ▶ **Show a control flow** involving signals that prints 19 twice
- ▶ **Why is `f()` not reentrant?**

```
1 int z;
2 int f(int x, int y){
3     int tmp = x + y;
4     z = tmp * 2 + 1;
5     return z;
6 }
7
8 void handle_signal(int sig){
9     int t = f(4,5);
10    printf("%d\n",t);
11    return;
12 }
13
14 int main(){
15     signal(SIGINT,handle_signal);
16     int v = f(1,2);
17     printf("%d\n",v);
18 }
```

Answer: Non-Reentrant Function Example

- ▶ Program below calls non-reentrant function `f()` in `main()` and `handle_signal()`
- ▶ With no interrupts, would expect to see 7 printed, with interrupts see 19 and 7
- ▶ Right hand shows one possible flow through the code which produces 19 then 19 again

```
1 int z;
2 int f(int x, int y){
3     int tmp = x + y;
4     z = tmp * 2 + 1;
5     return z;
6 }
7
8 void handle_signal(int sig){
9     int t = f(4,5);
10    printf("%d\n",t);
11    return;
12 }
13
14 int main(){
15     signal(SIGINT,handle_signal);
16     int v = f(1,2);
17     printf("%d\n",v);
18 }
```

```
EXECUTION STARTS IN main()
15: signal(SIGINT,handle_signal);
16: int v = f(1,2); // main(), Expect: (1+2)*2+1 = 7
    3: tmp = x + y;    // f(1,2): tmp = 1+2 = 3
    4: z = tmp*2 + 1;  // z is 7
SIGINT delivered, run handler
    9: int t = f(4,5); // handle_signal(2)
    3: tmp = x + y;    // f(4,5): tmp = 4+5 = 9
    4: z = tmp*2 + 1;  // z is now 19
    5: return z;       // back to handle_signal()
    9: int t = f(4,5); // finished, t is 19
   10: printf("%d\n",t); // puts 19 on screen
   11: return;         // back to normal control
    5: return z;       // back to main(), but z is 19
   16: int v = f(1,2);  // v is Actually 19
   17: printf("%d\n",v); // 19 Actually printed
                        // 7 Expected
```

Signal Sets

- ▶ A set of signals, likely implemented as a bit vector
- ▶ Functions allow addition, removal, clearing of set and tests for membership

```
#include <signal.h>

int sigemptyset(sigset_t *set);
// empty out the set

int sigfillset(sigset_t *set);
// fill the entire set with all signals

int sigaddset(sigset_t *set, int signo);
// add given signal to the set

int sigdelset(sigset_t *set, int signo);
// remove given signal to the set

// All of the above return 0 on succes, -1 on error

int sigismember(const sigset_t *set, int signo);
// return 1 if signal is a member of set, 0 if not
```

Examine sigsets-demo.c

Blocking (Disabling) Signals

- ▶ Processes can **block** signals, disable receiving them
- ▶ Signal is still there, just awaiting delivery
- ▶ Blocking is different from Ignoring a signal
 - ▶ Ignored signals are received and discarded
 - ▶ Blocked signals will be delivered after unblocking
- ▶ Can protect **Critical Sections** of code with by blocking if signals would screw it up

Process Signal Mask

Example: block all signals that can be blocked

```
sigset_t block_all, defaults;  
sigfillset( &block_all );  
sigprocmask(SIG_SETMASK, &block_all, &defaults);  
// contains all  
// block all signals  
// save defaults
```

Examine no-interruptions-block.c

Exercise: Protect Non-Reentrant Call

Examine the code for `non-reentrant.c` and modify it to use signal blocking to protect the **critical region** associated with calls to `getpwnam()`.

- ▶ Create a mask for all signals
- ▶ Block all signals prior to function call
- ▶ Unblock after returning
- ▶ Use code like below

```
sigset_t block_all, defaults;  
sigfillset( &block_all );  
sigprocmask(SIG_SETMASK, &block_all, &defaults);  
// contains all  
// block all signals  
// save defaults
```

Note: Be *very careful* where you unblock signal handling in `main()` to avoid errors: protect the **Critical Section**